


Generative AI COE

Technical Assessment
Job Candidate Generative AI Application





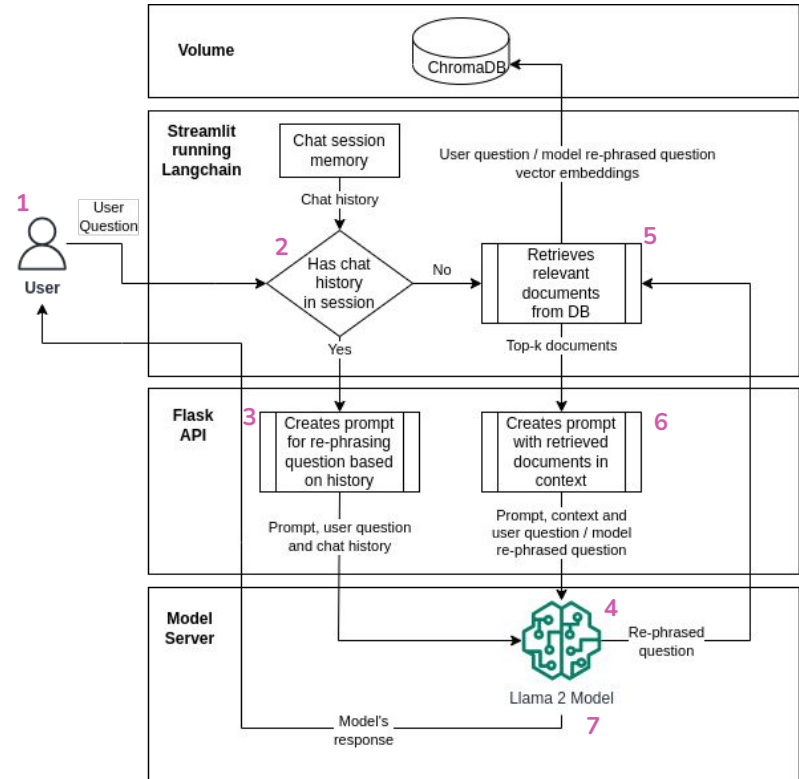
Demo



Click to be redirected to video in Google Drive

1.1 Application / Logic Flow Diagram

- 1, **User asks question** via Streamlit UI.
2. Langchain's conversational retrieval chain **checks** if there is any **chat history** in the same session.
3. If **yes, user question is augmented** to prompt the LLM to re-phrase the question based on the chat history.
4. **LLM re-phrases** the original user question, responds to Flask API which streams the response to Langchain.
5. Langchain **retrieves relevant documents** based on the original user question / model re-phrased question from the ChromaDB vector store.
6. **Final prompt** is created for LLM to respond to.
7. **LLM's response is streamed to the user** via Flask API and Streamlit application running Langchain





1.2 Application / Logic Flow Insights

1. How to chunk documents?

- All LLMs have limited input context length -> must chunk input documents and retrieve only most relevant documents.
- But chunking process can lead to loss of contextual information of chunks.
- Possible workarounds: Use metadata, summarise first then chunk or chunk depending on use-case (e.g. overlap, chunk size etc).

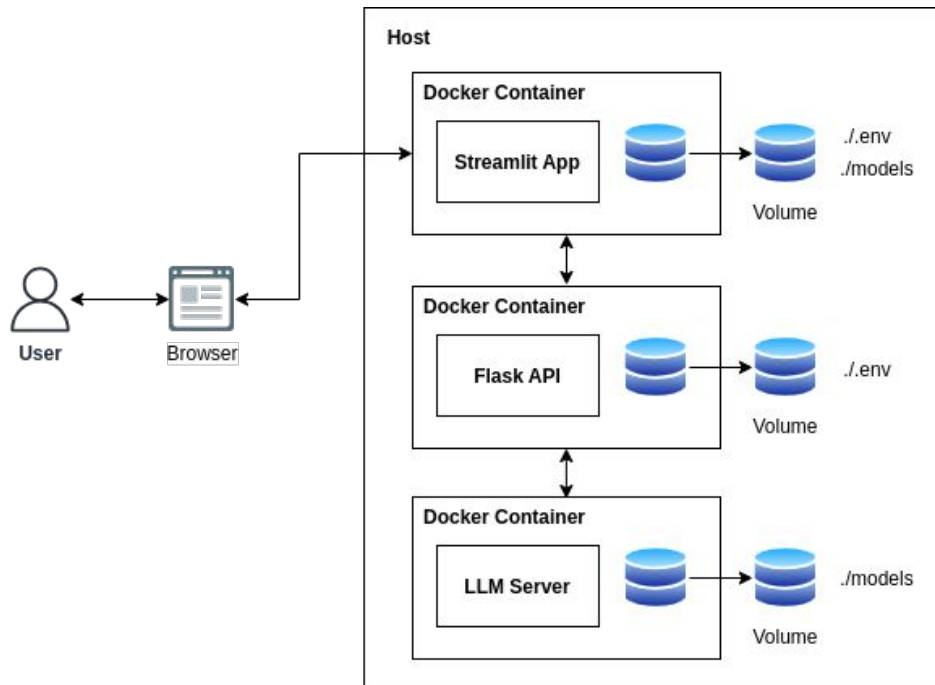
2. Which documents are more relevant?

- Some organisations have lots of documents with same information / conflicting information -> which is more accurate?
- Possible workarounds: Up-weight based on metadata (recency, author's seniority or user interactions), provide source document or cluster during retrieval.

3. Prompting

- Must engineer prompts: ensures that LLM does not answer irrelevant questions, LLM does not hallucinate and LLMs received the system prompt in the right format for model.

2.1 Architecture Diagram





2.2 Architecture Insights

Overall deployment considerations

- Differentiation of services by frontend, backend and LLM server based on purpose.
 - Frontend Streamlit: Serves UI to users and for defining the chat application interaction logic using Langchain.
 - Backend Flask API: Pre-processes user inputs into prompts specific to the task (e.g. chatbot for job candidates based on documents) and specific prompt templates used by the LLM.
 - Model Server: Hosts the LLM - LLama 2 Chat model.
 - **Advantages:**
 - Ability to change components easily. For example, change the model server to a different model, change the UI or change the database vector store independently.
 - Ability to add new components easily. For example, using different models for rephrasing questions and responding to final questions.
 - Ability to scale vertically or horizontally depending on the load requirements. For example, to use more powerful machines for model server.
 - Ability to re-purpose stack for other tasks. For example, instead of a chatbot for a job candidate can easily change the prompts in the API to do document answering for enterprise search.
 - Allows LLM server to be shared with other applications and not specifically for one application.

Thank you!

